

```
In [1]: # Load Packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import matplotlib.ticker as mtick
import sqlite3
import seaborn as sns
from imblearn.over_sampling import SMOTENC
from sklearn.linear_model import LinearRegression
from sklearn import tree, preprocessing
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, plot_confusion_matrix, recall_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier, RandomForestClassifier, ExtraTreesClassifier
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder, StandardScaler
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import roc_curve, auc, f1_score, make_scorer, recall_score
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from matplotlib.offsetbox import OffsetImage, AnnotationBbox
from matplotlib.pyplot import figure
from bs4 import BeautifulSoup
import time
import requests      # to get images
import shutil        # to save files locally
import datetime
from scipy.stats import norm
import warnings
warnings.filterwarnings('ignore')
import xgboost
from xgboost import XGBClassifier
from imblearn import under_sampling, over_sampling
from imblearn.over_sampling import SMOTE, ADASYN
import random
from random import randint
from sklearn.datasets import *
from IPython.display import Image, display_svg, SVG
import os
from sklearn.pipeline import Pipeline
from dtreeviz.trees import *
from sklearn.tree import plot_tree
os.environ["PATH"] += os.pathsep + "C:\\Users\\tmcro\\anaconda3\\pkgs\\graphviz-2.38-hfd603c8_2\\Library\\bin\\graphviz\\"
from sklearn.model_selection import train_test_split
from sklearn.compose import make_column_selector as selector
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from imblearn.pipeline import Pipeline as ImbPipeline # https://towardsdatascience.com/the-right-way-of-using-smote-with-cross-validation-92a8d09d00c7
```

```

from imblearn.pipeline import make_pipeline as make_imb_pipeline
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import make_pipeline
from sklearn.pipeline import FeatureUnion

```

```

In [2]: import pandas as pd
df = pd.read_csv('heart_2020_cleaned.csv')

target_name = "HeartDisease"
y = df[target_name]
X = df.drop(columns=[target_name])

```

```

In [3]: # Scoring Metric
def my_custom_score(y_true, y_pred):
    cf = confusion_matrix(y_true, y_pred)
    precision = cf[1,1] / sum(cf[:,1])
    recall = cf[1,1] / sum(cf[1,:])
    f1_score = 2*precision*recall / (precision + recall)
    rwf_score = (1+(2**2)) * ((precision * recall) / (((2**2) * precision)
+ recall))
    return rwf_score

my_scorer = make_scorer(my_custom_score, greater_is_better= True)

# Change class metric
class_metric = my_scorer

```

```

In [4]: X.dtypes

```

```

Out[4]: BMI                float64
Smoking                   object
AlcoholDrinking           object
Stroke                    object
PhysicalHealth            float64
MentalHealth              float64
DiffWalking              object
Sex                       object
AgeCategory               object
Race                     object
Diabetic                  object
PhysicalActivity          object
GenHealth                 object
SleepTime                 float64
Asthma                    object
KidneyDisease             object
SkinCancer                object
dtype: object

```

```

In [5]: numerical_columns_selector = selector(dtype_exclude=object)
categorical_columns_selector = selector(dtype_include=object)

numerical_columns = numerical_columns_selector(X)
categorical_columns = categorical_columns_selector(X)

```

```
In [6]: categori cal _col umns
```

```
Out[6]: [' Smoki ng' ,  
        ' Al cohol Dri nki ng' ,  
        ' Stroke' ,  
        ' Di ffWal ki ng' ,  
        ' Sex' ,  
        ' AgeCategory' ,  
        ' Race' ,  
        ' Di abeti c' ,  
        ' Physi cal Acti vi ty' ,  
        ' GenHeal th' ,  
        ' Asthma' ,  
        ' Ki dneyDi sease' ,  
        ' Ski nCancer' ]
```

```
In [7]: numeri cal _col umns
```

```
Out[7]: [' BMI ' , ' Physi cal Heal th' , ' Mental Heal th' , ' SleepTi me' ]
```

```
In [8]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

```
In [9]: catcols = [1, 2, 3, 6, 7, 8, 9, 10, 11, 12, 14, 15, 16]
```

```
In [10]: # To get the column names from onehotencoder  
ohe = OneHotEncoder(sparse=False, handle_unknown=' ignore' )  
checker = ohe. fi t_transform(X_train[categori cal _col umns])
```

```
In [11]: feature_names_categori cal = ohe.get_feature_names(categori cal _col umns)
```

```
In [12]: feature_names_categori cal
```

```
Out[12]: array([' Smoki ng_No' , ' Smoki ng_Yes' , ' Al cohol Dri nki ng_No' ,  
               ' Al cohol Dri nki ng_Yes' , ' Stroke_No' , ' Stroke_Yes' , ' Di ffWal ki ng_No' ,  
               ' Di ffWal ki ng_Yes' , ' Sex_Femal e' , ' Sex_Mal e' , ' AgeCategory_18-24' ,  
               ' AgeCategory_25-29' , ' AgeCategory_30-34' , ' AgeCategory_35-39' ,  
               ' AgeCategory_40-44' , ' AgeCategory_45-49' , ' AgeCategory_50-54' ,  
               ' AgeCategory_55-59' , ' AgeCategory_60-64' , ' AgeCategory_65-69' ,  
               ' AgeCategory_70-74' , ' AgeCategory_75-79' ,  
               ' AgeCategory_80 or ol der' , ' Race_Ameri can Indi an/Al askan Nat i ve' ,  
               ' Race_Asi an' , ' Race_Bl ack' , ' Race_Hi spani c' , ' Race_Other' ,  
               ' Race_Whi te' , ' Di abeti c_No' , ' Di abeti c_No, borderl i ne di abetes' ,  
               ' Di abeti c_Yes' , ' Di abeti c_Yes (duri ng pregnancy)' ,  
               ' Physi cal Acti vi ty_No' , ' Physi cal Acti vi ty_Yes' ,  
               ' GenHeal th_Excel l ent' , ' GenHeal th_Fai r' , ' GenHeal th_Good' ,  
               ' GenHeal th_Poor' , ' GenHeal th_Very good' , ' Asthma_No' , ' Asthma_Yes' ,  
               ' Ki dneyDi sease_No' , ' Ki dneyDi sease_Yes' , ' Ski nCancer_No' ,  
               ' Ski nCancer_Yes' ], dtype=object)
```

SMOTENC

```
In [13]: Smote_NC = SMOTENC(categorical_features= catcols, random_state=42)
```

```
In [14]: print(y_train.value_counts())  
# Fit SMOTE to training data  
X_train_resampled, y_train_resampled = Smote_NC.fit_resample(X_train, y_train)  
# Preview synthetic sample class distribution  
print('\n')  
print(pd.Series(y_train_resampled).value_counts())
```

```
No      219418  
Yes      20428  
Name: HeartDisease, dtype: int64
```

```
No      219418  
Yes      219418  
Name: HeartDisease, dtype: int64
```

```
In [15]: X_train = X_train_resampled  
y_train = y_train_resampled
```

```
In [16]: # Identify Processors  
categorical_preprocessor = OneHotEncoder(handle_unknown="ignore")  
numerical_preprocessor = StandardScaler()  
  
preprocessor = ColumnTransformer([('one-hot-encoder', categorical_preprocessor,  
                                  categorical_columns),  
                                  ('standard_scaler', numerical_preprocessor, numerical_columns)])  
  
smoter = SMOTENC(categorical_features= categorical_columns, random_state=0)  
classifier = LogisticRegression(max_iter=500)
```

Code Additions

```

In [17]: # SOURCE: The origin of this confusion matrix code was found on medium, '
# from https://medium.com/@dtuk81/confusion-matrix-visualization-fc31e3f30f
ea
def make_confusion_matrix(cf,
                           group_names=None,
                           categories='auto',
                           count=True,
                           percent=True,
                           cbar=True,
                           xyticks=True,
                           xyplotlabels=True,
                           sum_stats=True,
                           figsize=None,
                           cmap='Blues',
                           title=None):

    # CODE TO GENERATE SUMMARY STATISTICS & TEXT FOR SUMMARY STATS
    if sum_stats:
        #Accuracy is sum of diagonal divided by total observations
        accuracy = np.trace(cf) / float(np.sum(cf))

        #if it is a binary confusion matrix, show some more stats
        if len(cf)==2:
            #Metrics for Binary Confusion Matrices
            a = cf[0,0]
            b = cf[0,1]
            c = cf[1,0]
            d = cf[1,1]
            tn = ((a / (a+b))*100).round(2).astype(str) + '%'
            fp = ((b / (a+b))*100).round(2).astype(str) + '%'
            fn = ((c / (c+d))*100).round(2).astype(str) + '%'
            tp = ((d / (c+d))*100).round(2).astype(str) + '%'
            precision = cf[1,1] / sum(cf[:,1])
            recall = cf[1,1] / sum(cf[1,:])
            f1_score = 2*precision*recall / (precision + recall)
            rwf_score = (1+(2**2)) * ((precision * recall) / (((2**2) * pre
cision) + recall))
            stats_text = "\n\nAccuracy={: 0.3f}\nPrecision={: 0.3f}\nRecall =
{: 0.3f}\nF1 Score={: 0.3f}\n\nRecall-Weighted F Score={: 0.3f}".format(
                accuracy,precision,recall,f1_score,rwf_score)
        else:
            stats_text = "\n\nAccuracy={: 0.3f}".format(accuracy)
    else:
        stats_text = ""

    # CODE TO GENERATE TEXT INSIDE EACH SQUARE
    blanks = ['' for i in range(cf.size)]

    if group_names and len(group_names)==cf.size:
        group_labels = ["{}\n".format(value) for value in group_names]
    else:
        group_labels = blanks

    if count:
        group_counts = ["{0:0.0f}\n".format(value) for value in cf.flatten

```

```

()]
    else:
        group_counts = blanks

    if percent:
        group_percentages = [tn, fp, fn, tp]
        # old = group_percentages = ["{0:.2%}".format(value) for value in c
        f.flatten()/np.sum(cf)]
    else:
        group_percentages = blanks

    box_labels = [f"{v1}{v2}{v3}".strip() for v1, v2, v3 in zip(group_labels, group_counts, group_percentages)]
    box_labels = np.asarray(box_labels).reshape(cf.shape[0], cf.shape[1])

    # SET FIGURE PARAMETERS ACCORDING TO OTHER ARGUMENTS
    if figsize==None:
        #Get default figure size if not set
        figsize = plt.rcParams.get('figure.figsize')

    if xyticks==False:
        #Do not show categories if xyticks is False
        categories=False

    # MAKE THE HEATMAP VISUALIZATION
    plt.figure(figsize=figsize)
    sns.heatmap(cf, annot=box_labels, fmt="", cmap=cmap, cbar=cbar, xticklabels=categories, yticklabels=categories)

    if xyplotlabels:
        plt.ylabel('True Label', weight = 'bold')
        plt.xlabel('Predicted Label' + stats_text, weight = 'bold')
    else:
        plt.xlabel(stats_text)

    if title:
        plt.title(title, size = 20, weight = 'bold')

```

```

In [18]: dfcols = ['Model', 'RWF Score', 'F1', 'Recall', 'Precision', 'Accuracy']
model_summary = pd.DataFrame(columns=dfcols)
model_summary

```

Out[18]:

Model	RWF Score	F1	Recall	Precision	Accuracy
-------	-----------	----	--------	-----------	----------

```
In [19]: # Define Result Saving Initial Function
def save_result(cf, model_name):
    global model_summary
    accuracy = np.trace(cf) / float(np.sum(cf))
    precision = cf[1,1] / sum(cf[:,1])
    recall = cf[1,1] / sum(cf[1,:])
    f1_score = 2*precision*recall / (precision + recall)
    rwf_score = (1+(2**2)) * ((precision * recall) / (((2**2) * precision) + recall))
    row = [(model_name, rwf_score, f1_score, recall, precision, accuracy)]
    res = pd.DataFrame(columns = dfcols, data = row)
    yeep = [model_summary, res]
    model_summary = pd.concat(yeep)
    model_summary = model_summary.sort_values('RWF Score', ascending = False)
    model_summary = model_summary.drop_duplicates()
    return model_summary.round(3)
```

```
In [20]: def run_model(model, model_name):
    model.fit(X_train, y_train)
    model_prediction = model.predict(X_test)
    cf_matrix = confusion_matrix(y_test, model_prediction)
    save_result(cf_matrix, model_name)
    cf = make_confusion_matrix(cf_matrix)
    return model_summary
```

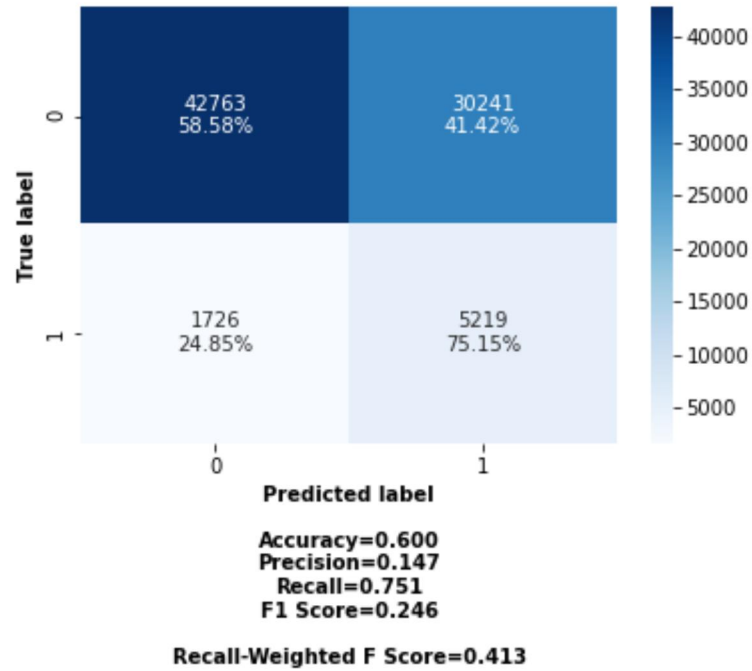
Models to Run

```
In [25]: # Initial Models
init_decision_tree = make_pipeline(preprocessor, DecisionTreeClassifier(random_state=0, max_depth=3))
log_reg = make_pipeline(preprocessor, LogisticRegression(random_state=0))
decision_tree = make_pipeline(preprocessor, DecisionTreeClassifier(max_depth=5, random_state=0))
random_forest = make_pipeline(preprocessor, RandomForestClassifier(n_estimators=50, random_state=0))
bagged_trees = make_pipeline(preprocessor, BaggingClassifier(DecisionTreeClassifier(), n_estimators=50, random_state=0))
extra_trees = make_pipeline(preprocessor, ExtraTreesClassifier(n_estimators=50, random_state=0))
knn = make_pipeline(preprocessor, KNeighborsClassifier(n_neighbors=5))
xgboost = make_pipeline(preprocessor, XGBClassifier(n_estimators=50, random_state=0))
```

```
In [26]: run_model (i n i t _ d e c i s i o n _ t r e e , ' I n i t i a l D e c i s i o n T r e e ' )
```

Out[26]:


	Model	RWF Score	F1	Recall	Precision	Accuracy
0	Logistic Regression	0.501597	0.337694	0.741541	0.218628	0.747326
0	Initial Decision Tree	0.412634	0.246150	0.751476	0.147180	0.600158



In [27]: `run_model(log_reg, 'Logistic Regression')`

Out[27]:


	Model	RWF Score	F1	Recall	Precision	Accuracy
0	Logistic Regression	0.501597	0.337694	0.741541	0.218628	0.747326
0	Initial Decision Tree	0.412634	0.246150	0.751476	0.147180	0.600158



```
In [28]: run_model (decisionTree, 'Decision Tree')
```

Out[28]:


	Model	RWF Score	F1	Recall	Precision	Accuracy
0	Logistic Regression	0.501597	0.337694	0.741541	0.218628	0.747326
0	Decision Tree	0.441529	0.276928	0.731317	0.170803	0.668251
0	Initial Decision Tree	0.412634	0.246150	0.751476	0.147180	0.600158



```
In [29]: run_model(random_forest, 'Random Forest')
```

Out[29]:


	Model	RWF Score	F1	Recall	Precision	Accuracy
0	Logistic Regression	0.501597	0.337694	0.741541	0.218628	0.747326
0	Decision Tree	0.441529	0.276928	0.731317	0.170803	0.668251
0	Initial Decision Tree	0.412634	0.246150	0.751476	0.147180	0.600158
0	Random Forest	0.304572	0.288472	0.316343	0.265114	0.864439



```
In [30]: run_model (baggedTrees, 'Bagged Trees')
```

Out[30]:


	Model	RWF Score	F1	Recall	Precision	Accuracy
0	Logistic Regression	0.501597	0.337694	0.741541	0.218628	0.747326
0	Decision Tree	0.441529	0.276928	0.731317	0.170803	0.668251
0	Initial Decision Tree	0.412634	0.246150	0.751476	0.147180	0.600158
0	Random Forest	0.304572	0.288472	0.316343	0.265114	0.864439
0	Bagged Trees	0.256132	0.252925	0.258315	0.247756	0.867440



```
In [31]: run_model (extraTrees, 'Extra Trees')
```

Out[31]:

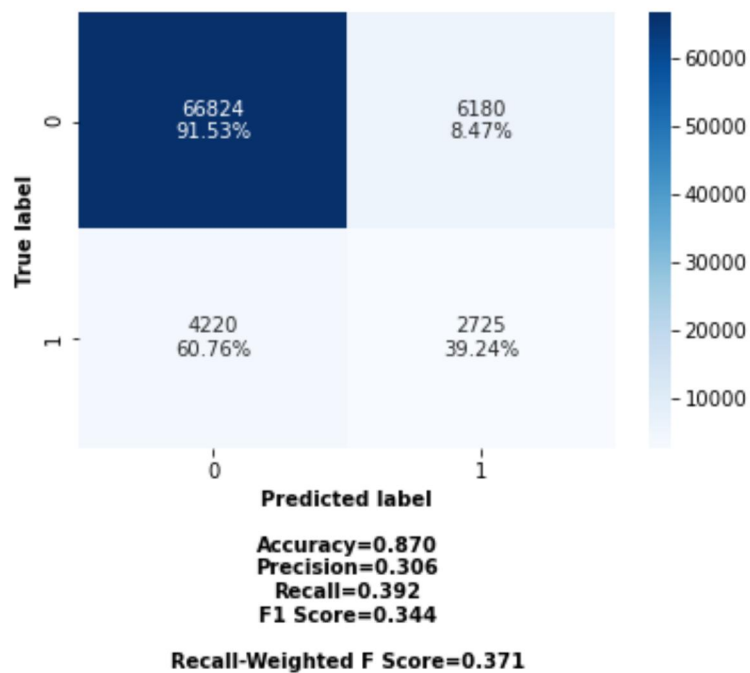
	Model	RWF Score	F1	Recall	Precision	Accuracy
0	Logistic Regression	0.501597	0.337694	0.741541	0.218628	0.747326
0	Decision Tree	0.441529	0.276928	0.731317	0.170803	0.668251
0	Initial Decision Tree	0.412634	0.246150	0.751476	0.147180	0.600158
0	Random Forest	0.304572	0.288472	0.316343	0.265114	0.864439
0	Extra Trees	0.300333	0.273155	0.321670	0.237357	0.851293
0	Bagged Trees	0.256132	0.252925	0.258315	0.247756	0.867440



```
In [32]: run_model(xgBoost, 'XGBoost')
```

Out[32]:

	Model	RWF Score	F1	Recall	Precision	Accuracy
0	Logistic Regression	0.501597	0.337694	0.741541	0.218628	0.747326
0	Decision Tree	0.441529	0.276928	0.731317	0.170803	0.668251
0	Initial Decision Tree	0.412634	0.246150	0.751476	0.147180	0.600158
0	XGBoost	0.371405	0.343849	0.392369	0.306008	0.869917
0	Random Forest	0.304572	0.288472	0.316343	0.265114	0.864439
0	Extra Trees	0.300333	0.273155	0.321670	0.237357	0.851293
0	Bagged Trees	0.256132	0.252925	0.258315	0.247756	0.867440



Iterate the best

```
In [33]: log_reg['logisticregression'].get_params()
```

```
Out[33]: {'C': 1.0,  
          'class_weight': None,  
          'dual': False,  
          'fit_intercept': True,  
          'intercept_scaling': 1,  
          'l1_ratio': None,  
          'max_iter': 100,  
          'multi_class': 'auto',  
          'n_jobs': None,  
          'penalty': 'l2',  
          'random_state': 0,  
          'solver': 'lbfgs',  
          'tol': 0.0001,  
          'verbose': 0,  
          'warm_start': False}
```

```
In [34]: # Define Grid  
rf_grid = [{'logisticregression__C': [1, 2],  
           'logisticregression__penalty': ['l2', 'none'],  
           'logisticregression__solver': ['newton-cg', 'lbfgs', 'sag']}]
```

```
In [35]: logreg_gridsearch = GridSearchCV(estimator=log_reg, param_grid=rf_grid, scoring=class_metric, cv=3)
```


```
In [36]: log_reg.get_params().keys()
```

```
Out[36]: dict_keys(['memory', 'steps', 'verbose', 'columntransformer', 'logisticregression', 'columntransformer__n_jobs', 'columntransformer__remainder', 'columntransformer__sparse_threshold', 'columntransformer__transformer_weights', 'columntransformer__transformers', 'columntransformer__verbose', 'columntransformer__one-hot-encoder', 'columntransformer__standard_scaler', 'columntransformer__one-hot-encoder__categories', 'columntransformer__one-hot-encoder__drop', 'columntransformer__one-hot-encoder__dtype', 'columntransformer__one-hot-encoder__handle_unknown', 'columntransformer__one-hot-encoder__sparse', 'columntransformer__standard_scaler__copy', 'columntransformer__standard_scaler__with_mean', 'columntransformer__standard_scaler__with_std', 'logisticregression__C', 'logisticregression__class_weight', 'logisticregression__dual', 'logisticregression__fit_intercept', 'logisticregression__intercept_scaling', 'logisticregression__l1_ratio', 'logisticregression__max_iter', 'logisticregression__multi_class', 'logisticregression__n_jobs', 'logisticregression__penalty', 'logisticregression__random_state', 'logisticregression__solver', 'logisticregression__tol', 'logisticregression__verbose', 'logisticregression__warm_start'])
```

```
In [37]: #logreg_gridsearch.fit(X_train, y_train)
run_model(logreg_gridsearch, 'Logistic Regression with Grid Search')
```

Out[37]:

	Model	RWF Score	F1	Recall	Precision	Accuracy
0	Logistic Regression	0.501597	0.337694	0.741541	0.218628	0.747326
0	Logistic Regression with Grid Search	0.501412	0.337529	0.741397	0.218502	0.747189
0	Decision Tree	0.441529	0.276928	0.731317	0.170803	0.668251
0	Initial Decision Tree	0.412634	0.246150	0.751476	0.147180	0.600158
0	XGBoost	0.371405	0.343849	0.392369	0.306008	0.869917
0	Random Forest	0.304572	0.288472	0.316343	0.265114	0.864439
0	Extra Trees	0.300333	0.273155	0.321670	0.237357	0.851293
0	Bagged Trees	0.256132	0.252925	0.258315	0.247756	0.867440



Winner: Logistic Regression


```
In [38]: log_reg.named_steps
```

```
Out[38]: {'columntransformer': ColumnTransformer(transformers=[('one-hot-encoder',  
OneHotEncoder(handle_unknown='ignore'),  
['Smoking', 'AlcoholDrinking', 'Stroke',  
'DiffWalking', 'Sex', 'AgeCategory', 'Race',  
'Diabetic', 'PhysicalActivity', 'GenHealth',  
'Asthma', 'KidneyDisease', 'SkinCancer']),  
('standard_scaler', StandardScaler(),  
['BMI', 'PhysicalHealth', 'MentalHealth',  
'SleepTime'])),  
'logisticregression': LogisticRegression(random_state=0)}
```

```
In [39]: coefs = log_reg['logisticregression'].coef_  
coefs = list(coefs[0])  
len(coefs)
```

```
Out[39]: 50
```

```
In [40]: numerical_columns
```

```
Out[40]: ['BMI', 'PhysicalHealth', 'MentalHealth', 'SleepTime']
```

```
In [41]: cat_list= list(feature_names_categorical)
```

```
In [42]: all_cats = cat_list + numerical_columns
```

In [43]: all_cats

```
Out[43]: ['Smoking_No',
'Smoking_Yes',
'AlcoholDrinking_No',
'AlcoholDrinking_Yes',
'Stroke_No',
'Stroke_Yes',
'DiffWalking_No',
'DiffWalking_Yes',
'Sex_Female',
'Sex_Male',
'AgeCategory_18-24',
'AgeCategory_25-29',
'AgeCategory_30-34',
'AgeCategory_35-39',
'AgeCategory_40-44',
'AgeCategory_45-49',
'AgeCategory_50-54',
'AgeCategory_55-59',
'AgeCategory_60-64',
'AgeCategory_65-69',
'AgeCategory_70-74',
'AgeCategory_75-79',
'AgeCategory_80 or older',
'Race_American Indian/Alaskan Native',
'Race_Asian',
'Race_Black',
'Race_Hispanic',
'Race_Other',
'Race_White',
'Diabetic_No',
'Diabetic_No, borderline diabetes',
'Diabetic_Yes',
'Diabetic_Yes (during pregnancy)',
'PhysicalActivity_No',
'PhysicalActivity_Yes',
'GenHealth_Excellent',
'GenHealth_Fair',
'GenHealth_Good',
'GenHealth_Poor',
'GenHealth_Very good',
'Asthma_No',
'Asthma_Yes',
'KidneyDisease_No',
'KidneyDisease_Yes',
'SkinCancer_No',
'SkinCancer_Yes',
'BMI',
'Physical Health',
'Mental Health',
'SleepTime']
```

```
In [44]: log_reg['logistic regression'].coef_
```

```
Out[44]: array([[ -4.05333675e-01,  -2.90920559e-02,   3.37922741e-01,
  -7.72348472e-01,  -6.18184131e-01,   1.83758400e-01,
  -3.03145729e-01,  -1.31280001e-01,  -6.00230712e-01,
    1.65804981e-01,  -2.16451588e+00,  -2.27540768e+00,
  -2.01627123e+00,  -1.44242048e+00,  -8.19559992e-01,
  -4.64338740e-01,   2.44673437e-01,   6.01561770e-01,
    1.01634869e+00,   1.28097544e+00,   1.61177442e+00,
    1.81754057e+00,   2.17521394e+00,  -6.26414042e-01,
  -7.51741889e-01,   6.32556488e-02,   3.01197732e-01,
  -4.49489479e-01,   1.02876630e+00,   5.26667578e-01,
  -4.80360117e-01,   1.00643036e+00,  -1.48716355e+00,
  -2.57080607e-01,  -1.77345124e-01,  -1.34099153e+00,
    6.31534785e-01,   2.00928726e-03,   9.89670893e-01,
  -7.16649164e-01,  -2.70863747e-01,  -1.63561984e-01,
  -2.54662829e-01,  -1.79762902e-01,  -1.43987311e-01,
  -2.90438420e-01,   9.58713097e-02,  -1.40068493e-02,
    9.92067232e-02,  -6.18929764e-02]])
```

```
In [45]: data1 = log_reg['logistic regression'].coef_  
cols = all_cats  
data = pd.DataFrame(data1, columns=cols)  
data = data.T  
data
```

Out[45]:

0

Smoking_No	-0.405334
Smoking_Yes	-0.029092
AlcoholDrinking_No	0.337923
AlcoholDrinking_Yes	-0.772348
Stroke_No	-0.618184
Stroke_Yes	0.183758
DiffWalking_No	-0.303146
DiffWalking_Yes	-0.131280
Sex_Female	-0.600231
Sex_Male	0.165805
AgeCategory_18-24	-2.164516
AgeCategory_25-29	-2.275408
AgeCategory_30-34	-2.016271
AgeCategory_35-39	-1.442420
AgeCategory_40-44	-0.819560
AgeCategory_45-49	-0.464339
AgeCategory_50-54	0.244673
AgeCategory_55-59	0.601562
AgeCategory_60-64	1.016349
AgeCategory_65-69	1.280975
AgeCategory_70-74	1.611774
AgeCategory_75-79	1.817541
AgeCategory_80 or older	2.175214
Race_American Indian/Alaskan Native	-0.626414
Race_Asian	-0.751742
Race_Black	0.063256
Race_Hispanic	0.301198
Race_Other	-0.449489
Race_White	1.028766
Diabetic_No	0.526668
Diabetic_No, borderline diabetes	-0.480360

0

Diabetic_Yes	1.006430
Diabetic_Yes (during pregnancy)	-1.487164
PhysicalActivity_No	-0.257081
PhysicalActivity_Yes	-0.177345
GenHealth_Excellent	-1.340992
GenHealth_Fair	0.631535
GenHealth_Good	0.002009
GenHealth_Poor	0.989671
GenHealth_Very good	-0.716649
Asthma_No	-0.270864
Asthma_Yes	-0.163562
KidneyDisease_No	-0.254663
KidneyDisease_Yes	-0.179763
SkinCancer_No	-0.143987

In [46]: data = data.reset_index()

```

In [55]: plt.figure(figsize=(16, 12))


conds = [(data[0] > 0),
         (data[0] < 0)]

vals = ['Positive', 'Negative']
data['pos_neg'] = np.select(conds, vals)
color_dict = {'Positive': '#49416D', 'Negative': '#E08D79'}

ax = sns.barplot(x=0, y='index', data=data, hue='pos_neg', palette=color_dict)
plt.ylabel('Categories', size=16, weight='bold')
plt.xlabel('Coefficient', size=16, weight='bold')
plt.title('Logistic Regression Coefficients', size=22, weight='bold')

```

Out[55]: Text(0.5, 1.0, 'Logistic Regression Coefficients')



In []: